

Unit 8: Handling data with pandas

Richard Foltyn

November 27, 2020

Contents

1	Handling data with pandas	1
1.1	Motivation	1
1.2	Creating pandas data structures	2
1.3	Viewing data	4
1.4	Indexing	5
1.5	Aggregation and reduction	10
1.6	Visualisation	12
2	Exercises	17
3	Solutions	17

1 Handling data with pandas

1.1 Motivation

So far, we have encountered NumPy arrays as the only way to store numerical data (we mostly ignored the built-in containers provided directly in Python).

However, while NumPy arrays are great for storing homogenous data which does not have any particular structure, they are somewhat limited when we want to use them for high-level data analysis.

For example, we usually want to process data sets with

1. several variables;
2. multiple observations, which need not be identical across variables (some values may be missing);
3. non-homogenous data types: for examples, names need to be stored as strings, birthdays as dates and income as a floating-point number.

While NumPy can in principle handle such situations, it puts all the burden on the user. Most users would prefer to not have to deal with such low-level details.

Imagine we want to store names, birth dates and annual income for two people:

Name	Date of birth	Income
Alice	1985-01-01	30,000
Bob	1997-05-12	-

No income was reported for Bob, so it's missing. With NumPy, we could do this as follows:

```
[1]: import numpy as np
      from datetime import date

      date1 = date(1985, 1, 1)           # birth date for Alice
      date2 = date(1997, 5, 12)         # birth date for Bob
```

```
data = np.array(['Alice', date1, 30000.0],
                ['Bob', date2, None])

data
```

```
[1]: array(['Alice', datetime.date(1985, 1, 1), 30000.0],
          ['Bob', datetime.date(1997, 5, 12), None], dtype=object)
```

```
[2]: data.dtype           # print array data type
```

```
[2]: dtype('O')
```

While we can create such arrays, they are almost useless for data analysis, in particular since everything is stored as a generic object.

- To be fair, NumPy offers an alternative array type called “[record](#)” or “[structured](#)” array which can handle fields of different data types.

However, the pandas library offers much more beyond that, so there is little reason to use structured arrays.

Pandas was created to offer more versatile data structures that are straightforward to use for storing, manipulating and analysing heterogeneous data:

1. Data is clearly organised in *variables* and *observations*, similar to econometrics programs such as Stata.
2. Each variable is permitted to have a different data type.
3. We can use *labels* to select observations, instead of having to use a linear numerical index as with NumPy.

We could, for example, index a data set using National Insurance Numbers.

4. Pandas offers many convenient data aggregation and reduction routines that can be applied to subsets of data.

For example, we can easily group observations by city and compute average incomes.

5. Pandas also offers many convenient data import / export functions that go beyond what’s in NumPy.

Should we be using pandas at all times, then? No!

- For low-level tasks where performance is essential, use NumPy.
- For homogenous data without any particular data structure, use NumPy.
- On the other hand, if data is heterogeneous, needs to be imported from an external data source and cleaned or transformed before performing computations, use pandas.

There are numerous tutorials on pandas on the internet, so we will keep this unit short and illustrate only the main concepts. Useful references to additional material include:

- The official [user guide](#).
- The official [pandas cheat sheet](#) which nicely illustrates the most frequently used operations.
- The official [API reference](#) with details on every pandas object and function.
- There are numerous tutorials (including videos) available on the internet. See [here](#) for a list.

1.2 Creating pandas data structures

Pandas has two main data structures:

1. *Series* represents observations of a single variable.
2. *DataFrame* is a container for several variables. You can think of each individual column of a *DataFrame* as a *Series*, and each row represents one observation.

The easiest way to create a Series or DataFrame is to create them from pre-existing data.

To access pandas data structures and routines, we need to import them first. The near-universal convention is to make pandas available using the name `pd`:

```
[3]: import pandas as pd
```

Examples:

We can create a DataFrame from a NumPy array:

```
[4]: import numpy as np
import pandas as pd          # universal convention: import using pd
from numpy.random import default_rng

# Draw normally distributed data
rng = default_rng(123)
data = rng.normal(size=(10,3))

# Define variable (or column) names
varnames = ['A', 'B', 'C']

# Create pandas DataFrame
pd.DataFrame(data, columns=varnames)
```

```
[4]:
```

	A	B	C
0	-0.989121	-0.367787	1.287925
1	0.193974	0.920231	0.577104
2	-0.636464	0.541952	-0.316595
3	-0.322389	0.097167	-1.525930
4	1.192166	-0.671090	1.000269
5	0.136321	1.532033	-0.659969
6	-0.311795	0.337769	-2.207471
7	0.827921	1.541630	1.126807
8	0.754770	-0.145978	1.281902
9	1.074031	0.392621	0.005114

This code creates a DataFrame of three variables called A, B and C with 10 observations each.

Alternatively, we can create a DataFrame from non-homogenous data as follows:

```
[5]: # Names (strings)
names = ['Alice', 'Bob']

# Birth dates (datetime objects)
bdates = pd.to_datetime(['1985-01-01', '1997-05-12'])

# Incomes (floats)
incomes = np.array([35000, np.nan])          # code missing income as NaN

# create DataFrame from dictionary
pd.DataFrame({'Name': names, 'Birthdate': bdates, 'Income': incomes})
```

```
[5]:
```

	Name	Birthdate	Income
0	Alice	1985-01-01	35000.0
1	Bob	1997-05-12	NaN

If data types differ across columns, as in the above example, it is often convenient to create the DataFrame by passing a dictionary as an argument. Each key represents a column name and each corresponding value contains the data for that variable.

1.3 Viewing data

With large data sets, you hardly ever want to print the entire `DataFrame`. Pandas by default limits the amount of data shown. You can use the `head()` and `tail()` methods to explicitly display a specific number of rows from the top or the end of a `DataFrame`.

To illustrate, we use a data set of a few UK universities that contains their name, number of students and budget in million pounds (both from Wikipedia), and their [Times Higher Education \(THE\) ranking](#).

We read in the data stored in the file `universities.csv` (from the `data/` folder) like this:

```
[6]: import pandas as pd

# relative path to CSV file
file = '../data/universities.csv'

# Load sample data set of UK universities
df = pd.read_csv(file, sep=';')
```

We can now display the first and last three rows:

```
[7]: df.head(3)          # show first three rows
```

```
[7]:
```

	Institution	Country	Founded	Students	Budget	Rank
0	University of Glasgow	Scotland	1451	30805	626.5	92
1	University of Edinburgh	Scotland	1583	34275	1102.0	30
2	University of St Andrews	Scotland	1413	8984	251.2	201

```
[8]: df.tail(3)          # show last three rows
```

```
[8]:
```

	Institution	Country	Founded	Students	Budget	\
20	University of Stirling	Scotland	1967	9548	113.3	
21	Queen's University Belfast	Northern Ireland	1810	18438	369.2	
22	Swansea University	Wales	1920	20620	NaN	

	Rank
20	301
21	200
22	251

To quickly compute some descriptive statistics for the *numerical* variables in the `DataFrame`, we use `describe()`:

```
[9]: df.describe()
```

```
[9]:
```

	Founded	Students	Budget	Rank
count	23.000000	23.000000	22.000000	23.000000
mean	1745.652174	24106.782609	768.609091	124.739130
std	256.992149	9093.000735	608.234948	104.498463
min	1096.000000	8984.000000	113.300000	1.000000
25%	1589.000000	18776.500000	340.850000	32.500000
50%	1826.000000	23247.000000	643.750000	107.000000
75%	1941.500000	30801.500000	1023.500000	195.500000
max	2004.000000	41180.000000	2450.000000	401.000000

Note that this automatically ignores the columns `Institution` and `Country` as they contain strings and computing the mean, etc. of a string variable does not make sense.

To see low-level information about the data type used in each column, we call `info()`:

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23 entries, 0 to 22
```

```
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Institution      23 non-null    object
1   Country           23 non-null    object
2   Founded           23 non-null    int64
3   Students          23 non-null    int64
4   Budget            22 non-null    float64
5   Rank              23 non-null    int64
dtypes: float64(1), int64(3), object(2)
memory usage: 1.2+ KB
```

Note that pandas automatically discards missing information in computations. For example, the budget for Swansea University is missing, and this observation is therefore dropped when computing descriptive statistics. Consequently, you see in the above table that the column Budget contains only 22 *non-null* entries.

1.4 Indexing

Pandas supports two types of indexing:

1. Indexing by position. This is basically identical to the indexing of other Python and NumPy containers.
2. Indexing by label, ie. by the values assigned to the row or column index. These labels need not be integers in increasing order, as is the case for NumPy.

We will see how to assign labels below.

Pandas indexing is performed either by using brackets `[]`, or by using `.loc[]` for label indexing, or `.iloc[]` for positional indexing.

Indexing via `[]` can be somewhat confusing:

- specifying `df['name']` returns the column name as a Series object.
- On the other hand, specifying a range such as `df[5:10]` returns the *rows* associated with the *positions* 5,...,9.

Examples:

```
[11]: import pandas as pd

# Load sample data set of UK universities
df = pd.read_csv('../data/universities.csv', sep=';')
df['Institution']          # select a single column
```

```
[11]: 0      University of Glasgow
1      University of Edinburgh
2      University of St Andrews
3      University of Aberdeen
4      University of Strathclyde
5      LSE
6      UCL
7      University of Cambridge
8      University of Oxford
9      University of Warwick
10     Imperial College London
11     King's College London
12     University of Manchester
13     University of Bristol
14     University of Birmingham
15     Queen Mary University of London
16     University of York
```

```

17         University of Nottingham
18         University of Dundee
19         Cardiff University
20         University of Stirling
21         Queen's University Belfast
22         Swansea University
Name: Institution, dtype: object

```

```
[12]: df[['Institution', 'Rank']]      # select multiple columns using a list
```

```
[12]:
```

	Institution	Rank
0	University of Glasgow	92
1	University of Edinburgh	30
2	University of St Andrews	201
3	University of Aberdeen	178
4	University of Strathclyde	401
5	LSE	27
6	UCL	16
7	University of Cambridge	6
8	University of Oxford	1
9	University of Warwick	77
10	Imperial College London	11
11	King's College London	35
12	University of Manchester	51
13	University of Bristol	91
14	University of Birmingham	107
15	Queen Mary University of London	110
16	University of York	133
17	University of Nottingham	158
18	University of Dundee	201
19	Cardiff University	191
20	University of Stirling	301
21	Queen's University Belfast	200
22	Swansea University	251

To return the rows at positions 1, 2 and 3 we use

```
[13]: df[1:4]
```

```
[13]:
```

	Institution	Country	Founded	Students	Budget	Rank
1	University of Edinburgh	Scotland	1583	34275	1102.0	30
2	University of St Andrews	Scotland	1413	8984	251.2	201
3	University of Aberdeen	Scotland	1495	14775	219.5	178

Pandas follows the Python convention that indices are 0-based, and the endpoint of a slice is not included.

1.4.1 Selection by label

Pandas uses *labels* to index and align data. These can be integer values starting at 0 with increments of 1 for each additional element, which is the default, but they need not be.

For example, we can replace the row index and use the Roman lower-case characters a, b, c, ... as labels instead of integers:

```
[14]: import pandas as pd
df = pd.read_csv('../data/universities.csv', sep=';')

# Create list of lower-case letters which has same
# length as the number of observations.
index = [chr(97+i) for i in range(len(df))]      # len(df) returns number of obs.
df.index = index
```

```
[14]: ['a',
      'b',
      'c',
      'd',
      'e',
      'f',
      'g',
      'h',
      'i',
      'j',
      'k',
      'l',
      'm',
      'n',
      'o',
      'p',
      'q',
      'r',
      's',
      't',
      'u',
      'v',
      'w']
```

```
[15]: df['index'] = index                                # create new column 'index'
      df.set_index(keys=['index'], inplace=True)         # set letters as index!

      # print first 3 rows using labels
      df['a':'c']                                         # This is the same as df[:3]
```

```
[15]:
```

	Institution	Country	Founded	Students	Budget	Rank
index						
a	University of Glasgow	Scotland	1451	30805	626.5	92
b	University of Edinburgh	Scotland	1583	34275	1102.0	30
c	University of St Andrews	Scotland	1413	8984	251.2	201

To add to the confusion, note that when specifying a range in terms of labels, the last element *is* included! Hence the row with index *c* in the above example is shown.

To more clearly distinguish between selection by label and by position, pandas provides the `.loc[]` and `.iloc[]` methods of indexing. To make your intention obvious, you should therefore adhere to the following rules:

1. Use `df['name']` only to select *columns* and nothing else.
2. Use `.loc[]` to select by label.
3. Use `.iloc[]` to select by position.

To illustrate, using `.loc[]` unambiguously indexes by label:

```
[16]: df.loc['d':'f', ['Institution', 'Students']]
```

```
[16]:
```

	Institution	Students
index		
d	University of Aberdeen	14775
e	University of Strathclyde	22640
f	LSE	11850

With `.loc[]` we can even perform slicing on column names, which is not possible with the simpler `df[]` syntax:

```
[17]: df.loc['d':'f', 'Institution':'Founded']
```

```
[17]:
```

	Institution	Country	Founded
d	University of Aberdeen	Scotland	1495
e	University of Strathclyde	Scotland	1964
f	LSE	England	1895

This includes all the columns between Institution and Founded, where the latter is again included since we are slicing by label.

Trying to pass in positional arguments will return an error if the index does not happen to be of type integer and the given values are actual labels:

```
[18]: df.loc[0:4]
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-1-11cc54301474> in <module>
----> 1 df.loc[0:4]

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexing.py in
-> _getitem__(self, key)
    877
    878         maybe_callable = com.apply_if_callable(key, self.obj)
--> 879         return self._getitem_axis(maybe_callable, axis=axis)
    880
    881     def _is_scalar_access(self, key: Tuple):

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexing.py in
-> _getitem_axis(self, key, axis)
   1086         if isinstance(key, slice):
   1087             self._validate_key(key, axis)
-> 1088             return self._get_slice_axis(key, axis=axis)
   1089         elif com.is_bool_indexer(key):
   1090             return self._getbool_axis(key, axis=axis)

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexing.py in
-> _get_slice_axis(self, slice_obj, axis)
   1121         labels = obj._get_axis(axis)
   1122         indexer = labels.slice_indexer(
-> 1123             slice_obj.start, slice_obj.stop, slice_obj.step, kind="loc"
   1124         )
   1125

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexes/base.py
-> in slice_indexer(self, start, end, step, kind)
   4964         slice(1, 3, None)
   4965         """
-> 4966         start_slice, end_slice = self.slice_locs(start, end, step=step,
-> kind=kind)
   4967
   4968         # return a slice

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexes/base.py
-> in slice_locs(self, start, end, step, kind)
   5165         start_slice = None
   5166         if start is not None:
-> 5167             start_slice = self.get_slice_bound(start, "left", kind)
   5168         if start_slice is None:
   5169             start_slice = 0

~/.conda/envs/py3-default/lib/python3.7/site-packages/pandas/core/indexes/base.py
-> in get_slice_bound(self, label, side, kind)
```



```

5077         # For datetime indices label may be a string that has to be
↳converted
5078         # to datetime boundary according to its resolution.
-> 5079         label = self._maybe_cast_slice_bound(label, side, kind)
5080
5081         # we need to look up the label

~/.../python3.7/site-packages/pandas/core/indexes/base.py
↳in _maybe_cast_slice_bound(self, label, side, kind)
5029         # this is rejected (generally .loc gets you here)
5030         elif is_integer(label):
-> 5031             self._invalid_indexer("slice", label)
5032
5033         return label

~/.../python3.7/site-packages/pandas/core/indexes/base.py
↳in _invalid_indexer(self, form, key)
3266         """
3267         raise TypeError(
-> 3268             f"cannot do {form} indexing on {type(self).__name__} with
↳these "
3269             f"indexers [{key}] of type {type(key).__name__}"
3270         )

TypeError: cannot do slice indexing on Index with these indexers [0] of type int

```

Somewhat surprisingly, we can also pass boolean arrays to `.loc[]` even though these are clearly not labels:

```
[19]: df.loc[df['Country'] == 'Scotland']
```

```
[19]:
```

	Institution	Country	Founded	Students	Budget	Rank
index						
a	University of Glasgow	Scotland	1451	30805	626.5	92
b	University of Edinburgh	Scotland	1583	34275	1102.0	30
c	University of St Andrews	Scotland	1413	8984	251.2	201
d	University of Aberdeen	Scotland	1495	14775	219.5	178
e	University of Strathclyde	Scotland	1964	22640	304.4	401
s	University of Dundee	Scotland	1967	15915	256.4	201
u	University of Stirling	Scotland	1967	9548	113.3	301

Indexing via `.loc[]` supports a few more types of arguments, see the [official documentation](#) for details.

1.4.2 Selection by position

Conversely, if we want to select items exclusively by their position and ignore their labels, we use `.iloc[]`:

```
[20]: df.iloc[0:4, 0:2] # select first 4 rows, first 2 columns
```

```
[20]:
```

	Institution	Country
index		
a	University of Glasgow	Scotland
b	University of Edinburgh	Scotland
c	University of St Andrews	Scotland
d	University of Aberdeen	Scotland

Again, `.iloc[]` supports a multitude of other arguments, including boolean arrays. See the [official documentation](#) for details.

1.5 Aggregation and reduction

1.5.1 Working with entire DataFrames

The simplest way to perform data reduction is to invoke the desired routine on the entire DataFrame:

```
[21]: import pandas as pd

df = pd.read_csv('../data/universities.csv', sep=';')
df.mean()
```

```
[21]: Founded      1745.652174
      Students    24106.782609
      Budget      768.609091
      Rank        124.739130
      dtype: float64
```

Methods such as `mean()` are by default applied column-wise to each numerical column.

One big advantage over NumPy is that missing values (represented by `np.nan`) are automatically ignored:

```
[22]: # Set Aberdeen's THE ranking to missing
df.loc[df['Institution'] == 'University of Aberdeen', 'Rank'] = np.nan
df.loc[df['Institution'] == 'University of Aberdeen']
```

```
[22]:
```

	Institution	Country	Founded	Students	Budget	Rank
3	University of Aberdeen	Scotland	1495	14775	219.5	NaN

```
[23]: # mean() automatically drops missing observations
df['Rank'].mean()
```

```
[23]: 122.31818181818181
```

1.5.2 Splitting and grouping

Applying aggregation functions to the entire DataFrame is similar to what we can do with NumPy. The added flexibility of pandas becomes obvious once we want to apply these functions to subsets of data, ie. groups, which we can define based on values or index labels.

For example, we can easily group our universities by country:

```
[24]: import pandas as pd

df = pd.read_csv('../data/universities.csv', sep=';')

groups = df.groupby(['Country'])
```

Here `groups` is a special pandas objects which can subsequently be used to process group-specific data. To compute the group-wise averages, we can simply run

```
[25]: groups.mean()
```

```
[25]:
```

	Founded	Students	Budget	Rank
Country				
England	1745.923077	27119.846154	1001.700000	63.307692
Northern Ireland	1810.000000	18438.000000	369.200000	200.000000
Scotland	1691.428571	19563.142857	410.471429	200.571429
Wales	1901.500000	23259.000000	644.800000	221.000000

Groups support column indexing: if we want to only compute the total number of students for each country in our sample, we can do this as follows:

```
[26]: groups['Students'].sum()
```

```
[26]: Country
      England          352558
      Northern Ireland    18438
      Scotland          136942
      Wales              46518
      Name: Students, dtype: int64
```

There are numerous routines to aggregate grouped data, for example:

- `mean()`, `sum()`: averages and sums over numerical items within groups.
- `std()`, `var()`: within-group std. dev. and variances
- `size()`: group sizes
- `first()`, `last()`: first and last elements in each group
- `min()`, `max()`: minimum and maximum elements within a group

Examples:

```
[27]: groups.size()          # return number of elements in each group
```

```
[27]: Country
      England          13
      Northern Ireland    1
      Scotland           7
      Wales              2
      dtype: int64
```

```
[28]: groups.first()        # return first element in each group
```

```
[28]:
```

	Institution	Founded	Students	Budget	Rank
Country					
England	LSE	1895	11850	415.1	27
Northern Ireland	Queen's University Belfast	1810	18438	369.2	200
Scotland	University of Glasgow	1451	30805	626.5	92
Wales	Cardiff University	1883	25898	644.8	191

We can create custom aggregation routines by calling `agg()` or `aggregate()` on the grouped object. To illustrate, we count the number of universities in each country that have more than 20,000 students:

```
[29]: groups['Students'].agg(lambda x: np.sum(x >= 20000))
```

```
[29]: Country
      England          10
      Northern Ireland    0
      Scotland           3
      Wales              2
      Name: Students, dtype: int64
```

Note that we called `agg()` only on the column `Students`, otherwise the function would be applied to every column separately, which is not what we want.

The most flexible aggregation method is `apply()` which calls a given function, passing the entire group-specific subset of data (including all columns) as an argument, and glues together the results.

For example, if we want to compute the average budget per student (in pounds), we can do this as follows:

```
[30]: # Budget is in millions of pounds, rescale by 1.0e6
      groups.apply(lambda x: x['Budget'].sum() / x['Students'].sum() * 1.0e6)
```

```
[30]: Country
      England          36936.050239
```

```
Northern Ireland    20023.863760
Scotland            20981.875539
Wales               13861.301002
dtype: float64
```

We couldn't have done this with `agg()`, since `agg()` never gets to see the entire chunk of data but only one column at a time.

This section provided only a first look at pandas's "split-apply-combine" functionality implemented via `groupby`. See the [official documentation](#) for more details.

1.6 Visualisation

We covered plotting with Matplotlib in earlier units. Pandas itself implements some convenience wrappers around Matplotlib plotting routines which allow us to quickly inspect data stored in `DataFrames`. Alternatively, we can extract the numerical data and pass it to Matplotlib's routines manually.


For example, to plot student numbers as a bar chart, we can directly use pandas:

```
[31]: import pandas as pd

df = pd.read_csv('../data/universities.csv', sep=';')

# set institution as label so they automatically show up in plot
df2 = df.set_index(keys=['Institution'])
df2['Students'].plot(kind='bar')
```


```
[31]: <AxesSubplot:xlabel='Institution'>
```



Alternatively, we can construct the graph using Matplotlib ourselves:

```
[32]: import matplotlib.pyplot as plt

labels = df['Institution'].to_list()          # labels as list
values = df['Students'].to_numpy()           # data as NumPy array
plt.bar(labels, values)
plt.tick_params(axis='x', labelrotation=90)
```



Sometimes Matplotlib's routines directly work with pandas's data structures, sometimes they don't. In cases where they don't, we can convert a `DataFrame` or `Series` object to a NumPy array using the `to_numpy()` method, and convert a `Series` to a Python list using `to_list()`, as illustrated in the example above.

To plot timeseries-like data, we can use the `plot()` method, which optionally accepts arguments to specify which columns should be used for the x-axis and which for the y-axis:


```
[33]: import numpy as np
import pandas as pd

# Instantiate RNG
rng = np.random.default_rng(123)

# Create pandas DataFrame
nobs = 50
df = pd.DataFrame({'A': np.arange(nobs), 'B': rng.normal(size=nobs)})

df.plot(x='A', y='B')           # plot A on x-axis, B on y-axis
```

```
[33]: <AxesSubplot:xlabel='A'>
```



To quickly generate some descriptive statistics, we can use the built-in box plot:


```
[34]: import numpy as np
import pandas as pd

# Instantiate RNG
rng = np.random.default_rng(123)

# Create pandas DataFrame
df = pd.DataFrame(rng.normal(size=(100, 5)),
                  columns=['A', 'B', 'C', 'D', 'E'])

df.plot.box()           # same as df.plot(kind='box')
```

```
[34]: <AxesSubplot:>
```




Similarly, we can generate scatter plots, plotting one column against another:

```
[35]: # Create pandas DataFrame
df = pd.DataFrame(rng.uniform(size=(100, 5)),
                  columns=['A', 'B', 'C', 'D', 'E'])


df.plot.scatter(x='A', y='E', label='Pandas scatter plot')
```

```
[35]: <AxesSubplot:xlabel='A', ylabel='E'>
```




```
[36]: # We can even use a column to specify the dot size!  
df.plot.scatter(x='A', y='E', s=df['B']*100.0)
```

```
[36]: <AxesSubplot:xlabel='A', ylabel='E'>
```



In general, the wrappers implemented in pandas are useful to get an idea how the data looks like. For reusable code or more complex graphs, we'll usually want to directly use Matplotlib and pass the data converted to NumPy arrays.

2 Exercises

TBA

3 Solutions

TBA